

# EE658/758 Machine Learning in Engineering Assignment

## 1. Data Preprocessing

```
In [2]: # Libraries
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split # For Data Splitting
df = pd.read_csv('insurance.csv')
df.head()
```

```
Out[2]:
```

	Age	Gender	BMI	Children	Smoker	Region	Expenses
0	19	female	27.9	0	yes	southwest	16884.92
1	18	male	33.8	1	no	southeast	1725.55
2	28	male	33.0	3	no	southeast	4449.46
3	33	male	22.7	0	no	northwest	21984.47
4	32	male	28.9	0	no	northwest	3866.86

```
In [3]: df.shape
```

```
Out[3]: (1338, 7)
```

```
In [4]: df.dtypes
```

```
Out[4]: Age          int64
Gender         object
BMI            float64
Children       int64
Smoker         object
Region         object
Expenses       float64
dtype: object
```

```
In [5]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1338 entries, 0 to 1337
Data columns (total 7 columns):
#   Column      Non-Null Count  Dtype
---  -
0   Age         1338 non-null   int64
1   Gender      1332 non-null   object
2   BMI         1330 non-null   float64
3   Children    1338 non-null   int64
4   Smoker      1338 non-null   object
5   Region      1336 non-null   object
6   Expenses    1337 non-null   float64
dtypes: float64(2), int64(2), object(3)
memory usage: 73.3+ KB
```

```
In [6]: df.describe()
```

Out[6]:

	Age	BMI	Children	Expenses
<b>count</b>	1338.000000	1330.000000	1338.000000	1337.000000
<b>mean</b>	39.207025	30.676917	1.136024	13273.306111
<b>std</b>	14.049960	6.094868	3.194662	12114.083012
<b>min</b>	18.000000	16.000000	-65.000000	1121.870000
<b>25%</b>	27.000000	26.300000	0.000000	4738.270000
<b>50%</b>	39.000000	30.400000	1.000000	9377.900000
<b>75%</b>	51.000000	34.700000	2.000000	16657.720000
<b>max</b>	64.000000	53.100000	70.000000	63770.430000

In [7]: `df.isnull().sum()`

Out[7]:

Age	0
Gender	6
BMI	8
Children	0
Smoker	0
Region	2
Expenses	1

dtype: int64

In [8]: `df[df['Gender'].isnull()]`

Out[8]:

	Age	Gender	BMI	Children	Smoker	Region	Expenses
<b>425</b>	45	NaN	24.3	5	no	southeast	9788.87
<b>572</b>	30	NaN	43.1	2	no	southeast	4753.64
<b>729</b>	41	NaN	36.1	1	no	southeast	6781.35
<b>914</b>	33	NaN	24.6	2	no	northwest	5257.51
<b>1313</b>	19	NaN	34.7	2	yes	southwest	36397.58
<b>1334</b>	18	NaN	31.9	0	no	northeast	2205.98

In [9]: `df[df['BMI'].isnull()]`

Out[9]:

	Age	Gender	BMI	Children	Smoker	Region	Expenses
<b>8</b>	37	male	NaN	2	no	northeast	6406.41
<b>283</b>	55	female	NaN	1	no	northeast	11879.10
<b>580</b>	59	male	NaN	1	no	northeast	12913.99
<b>769</b>	38	female	NaN	2	no	northwest	6933.24
<b>946</b>	42	male	NaN	2	no	southwest	7160.09
<b>1100</b>	33	female	NaN	2	yes	northeast	16776.30
<b>1235</b>	26	male	NaN	0	no	northwest	2699.57
<b>1335</b>	18	female	NaN	0	no	southeast	1629.83

In [10]: `df[df['Expenses'].isnull()]`

Out[10]:

	Age	Gender	BMI	Children	Smoker	Region	Expenses
<b>810</b>	46	female	30.8	3	no	southwest	NaN

```
In [11]: df[df['Region'].isnull()]
```

```
Out[11]:
```

	Age	Gender	BMI	Children	Smoker	Region	Expenses
652	48	female	31.1	0	no	NaN	8280.62
788	29	male	22.5	3	no	NaN	5209.58

```
In [12]: df.dropna(inplace=True)
```

```
In [13]: df.isnull().sum()
```

```
Out[13]: Age      0
Gender    0
BMI       0
Children  0
Smoker    0
Region    0
Expenses  0
dtype: int64
```

```
In [14]: df[df.duplicated()]
```

```
Out[14]:
```

	Age	Gender	BMI	Children	Smoker	Region	Expenses
581	19	male	30.6	0	no	northwest	1639.56

```
In [15]: df[(df['Age'] == 19) & (df['Gender'] == 'male') & (df['BMI'] == 30.6)]
```

```
Out[15]:
```

	Age	Gender	BMI	Children	Smoker	Region	Expenses
195	19	male	30.6	0	no	northwest	1639.56
581	19	male	30.6	0	no	northwest	1639.56

```
In [16]: df.drop_duplicates(inplace=True)
```

```
In [17]: df[df.duplicated()]
```

```
Out[17]:
```

Age	Gender	BMI	Children	Smoker	Region	Expenses
-----	--------	-----	----------	--------	--------	----------

```
In [18]: df['Gender'].value_counts()
```

```
Out[18]: Gender
male      668
female    652
Name: count, dtype: int64
```

```
In [19]: df['Smoker'].value_counts()
```

```
Out[19]: Smoker
no       1048
yes       272
Name: count, dtype: int64
```

```
In [20]: df['Region'].value_counts()
```

```
Out[20]: Region
southeast    359
southwest    322
northwest    321
northeast    318
Name: count, dtype: int64
```

## Encoding

```
In [21]: gender_mapping = {'male': 0, 'female': 1}
smoker_mapping = {'no': 0, 'yes': 1}

df['Gender'] = df['Gender'].map(gender_mapping)
df['Smoker'] = df['Smoker'].map(smoker_mapping)
```

```
In [22]: df.head()
```

```
Out[22]:
```

	Age	Gender	BMI	Children	Smoker	Region	Expenses
0	19	1	27.9	0	1	southwest	16884.92
1	18	0	33.8	1	0	southeast	1725.55
2	28	0	33.0	3	0	southeast	4449.46
3	33	0	22.7	0	0	northwest	21984.47
4	32	0	28.9	0	0	northwest	3866.86

```
In [23]: region_mapping = {'southeast': 1, 'southwest': 2, 'northwest': 3, 'northeast': 4}
df['Region'] = df['Region'].map(region_mapping)
```

```
In [24]: df.head()
```

```
Out[24]:
```

	Age	Gender	BMI	Children	Smoker	Region	Expenses
0	19	1	27.9	0	1	2	16884.92
1	18	0	33.8	1	0	1	1725.55
2	28	0	33.0	3	0	1	4449.46
3	33	0	22.7	0	0	3	21984.47
4	32	0	28.9	0	0	3	3866.86

```
In [25]: df.dtypes
```

```
Out[25]: Age          int64
Gender          int64
BMI            float64
Children        int64
Smoker          int64
Region          int64
Expenses       float64
dtype: object
```

## Min Max Scaling

```
In [26]: from sklearn.preprocessing import MinMaxScaler  
scaler = MinMaxScaler()
```

```
In [27]: # Expenses column is Target Column  
features_to_normalize = ['Age', 'Gender', 'BMI', 'Children', 'Smoker', 'Region']  
features_to_normalize
```

```
Out[27]: ['Age', 'Gender', 'BMI', 'Children', 'Smoker', 'Region']
```

```
In [28]: df[features_to_normalize] = scaler.fit_transform(df[features_to_normalize])
```

```
In [29]: df.head()
```

```
Out[29]:
```

	Age	Gender	BMI	Children	Smoker	Region	Expenses
0	0.021739	1.0	0.320755	0.481481	1.0	0.333333	16884.92
1	0.000000	0.0	0.479784	0.488889	0.0	0.000000	1725.55
2	0.217391	0.0	0.458221	0.503704	0.0	0.000000	4449.46
3	0.326087	0.0	0.180593	0.481481	0.0	0.666667	21984.47
4	0.304348	0.0	0.347709	0.481481	0.0	0.666667	3866.86

## 2. Splitting the Data

```
In [30]: X = df.drop('Expenses',axis=1) # Features  
Y = df['Expenses'] # Labels
```

```
In [31]: X.head()
```



```
Out[31]:
```

	Age	Gender	BMI	Children	Smoker	Region
0	0.021739	1.0	0.320755	0.481481	1.0	0.333333
1	0.000000	0.0	0.479784	0.488889	0.0	0.000000
2	0.217391	0.0	0.458221	0.503704	0.0	0.000000
3	0.326087	0.0	0.180593	0.481481	0.0	0.666667
4	0.304348	0.0	0.347709	0.481481	0.0	0.666667

```
In [32]: Y.head()
```

```
Out[32]: 0    16884.92
1     1725.55
2     4449.46
3     21984.47
4      3866.86
Name: Expenses, dtype: float64
```

```
In [33]: # Split Ratio ---> 75:25
# Training : Testing
# Random_State for reproducibility
X_train, X_test, y_train, y_test = train_test_split(X, Y, test_size=0.25, random_state=42)
```

```
In [34]: print("X_train:\n", X_train.head(1))
print("\nX_test:\n", X_test.head(1))
print("\ny_train:\n", y_train.head(1))
print("\ny_test:\n", y_test.head(1))
```

```
X_train:
      Age  Gender      BMI  Children  Smoker   Region
587  0.347826    1.0  0.382749  0.488889    1.0  0.666667
```

```
X_test:
      Age  Gender      BMI  Children  Smoker   Region
685  0.76087    0.0  0.280323  0.496296    0.0    1.0
```

```
y_train:
587    43943.88
Name: Expenses, dtype: float64
```

```
y_test:
685    11244.38
Name: Expenses, dtype: float64
```

### 3. Gradient Descent Implementation

```
In [35]: X_train.head()
```

```
Out[35]:
```

	Age	Gender	BMI	Children	Smoker	Region
<b>587</b>	0.347826	1.0	0.382749	0.488889	1.0	0.666667
<b>356</b>	0.608696	0.0	0.752022	0.503704	0.0	0.000000
<b>992</b>	0.695652	1.0	0.420485	0.496296	0.0	0.333333
<b>464</b>	0.021739	0.0	0.247978	0.481481	0.0	0.666667
<b>950</b>	0.847826	0.0	0.061995	0.481481	0.0	1.000000

```
In [36]: y_train.head()
```

```
Out[36]: 587    43943.88
          356    8944.12
          992   10118.42
          464    1632.04
          950   11534.87
          Name: Expenses, dtype: float64
```

```
In [37]: from sklearn.metrics import mean_squared_error
def gradient_descent(X_train, y_train, learning_rate=0.01, n_iterations=10000, decay_rate=None):
    if len(X_train.shape) == 1: # For 1D Conversion
        X_train = X_train.reshape(-1, 1)

    if len(y_train.shape) == 1: # For 1D Conversion since ytrain is series so using numpy here
        y_train = y_train.to_numpy().reshape(-1, 1)

    # X bias term
    X_b = np.c_[np.ones((X_train.shape[0], 1)), X_train]

    # Data points and features
    m, n = X_train.shape[0], X_train.shape[1]

    # Weight matrix (W^T)
    W = np.random.randn(X_train.shape[1] + 1, 1)

    loss_values = []

    for iteration in range(n_iterations):
        gradients = 1/m * X_b.T.dot(X_b.dot(W) - y_train)

        if decay_rate:
            current_learning_rate = learning_rate / (1 + decay_rate * iteration)
        else:
            current_learning_rate = learning_rate

        W = W - current_learning_rate * gradients
        predictions = X_b.dot(W)
        loss_values.append(mean_squared_error(y_train, predictions))

    #result_string = f"Coefficient and Intercept of the Model: {W}\nLoss Values:\n{loss_values}"
    #print(result_string)
```

```
gradient_descent(X_train, y_train) # With Constant_Rate
```

```
In [38]: # Learning Rate Analysis with 0.01
gradient_descent(X_train, y_train, decay_rate=0.01)
```

```
In [39]: # Plot for Consant & Decay Learning Rates
def gradient_descent(X_train, y_train, learning_rate=0.01, n_iterations=10000, decay_rate=None):
    if len(X_train.shape) == 1: # For 1D Conversion
        X_train = X_train.reshape(-1, 1)

    if len(y_train.shape) == 1: # For 1D Conversion since ytrain is series so using numpy here
        y_train = y_train.to_numpy().reshape(-1, 1)

    # X bias term
    X_b = np.c_[np.ones((X_train.shape[0], 1)), X_train]

    # Data points and features
    m, n = X_train.shape[0], X_train.shape[1]

    # Weight matrix ( $W^T$ )
    W_constant = np.random.randn(n + 1, 1)
    W_decaying = np.random.randn(n + 1, 1)

    loss_values_constant = []
    loss_values_decaying = []

    for iteration in range(n_iterations):

        gradients_constant = 1/m * X_b.T.dot(X_b.dot(W_constant) - y_train)
        W_constant = W_constant - learning_rate * gradients_constant
        predictions_constant = X_b.dot(W_constant)
        loss_values_constant.append(mean_squared_error(y_train, predictions_constant))

        gradients_decaying = 1/m * X_b.T.dot(X_b.dot(W_decaying) - y_train)
        if decay_rate:
            current_learning_rate_decaying = learning_rate / (1 + decay_rate * iteration)
        else:
            current_learning_rate_decaying = learning_rate
        W_decaying = W_decaying - current_learning_rate_decaying * gradients_decaying
```

```

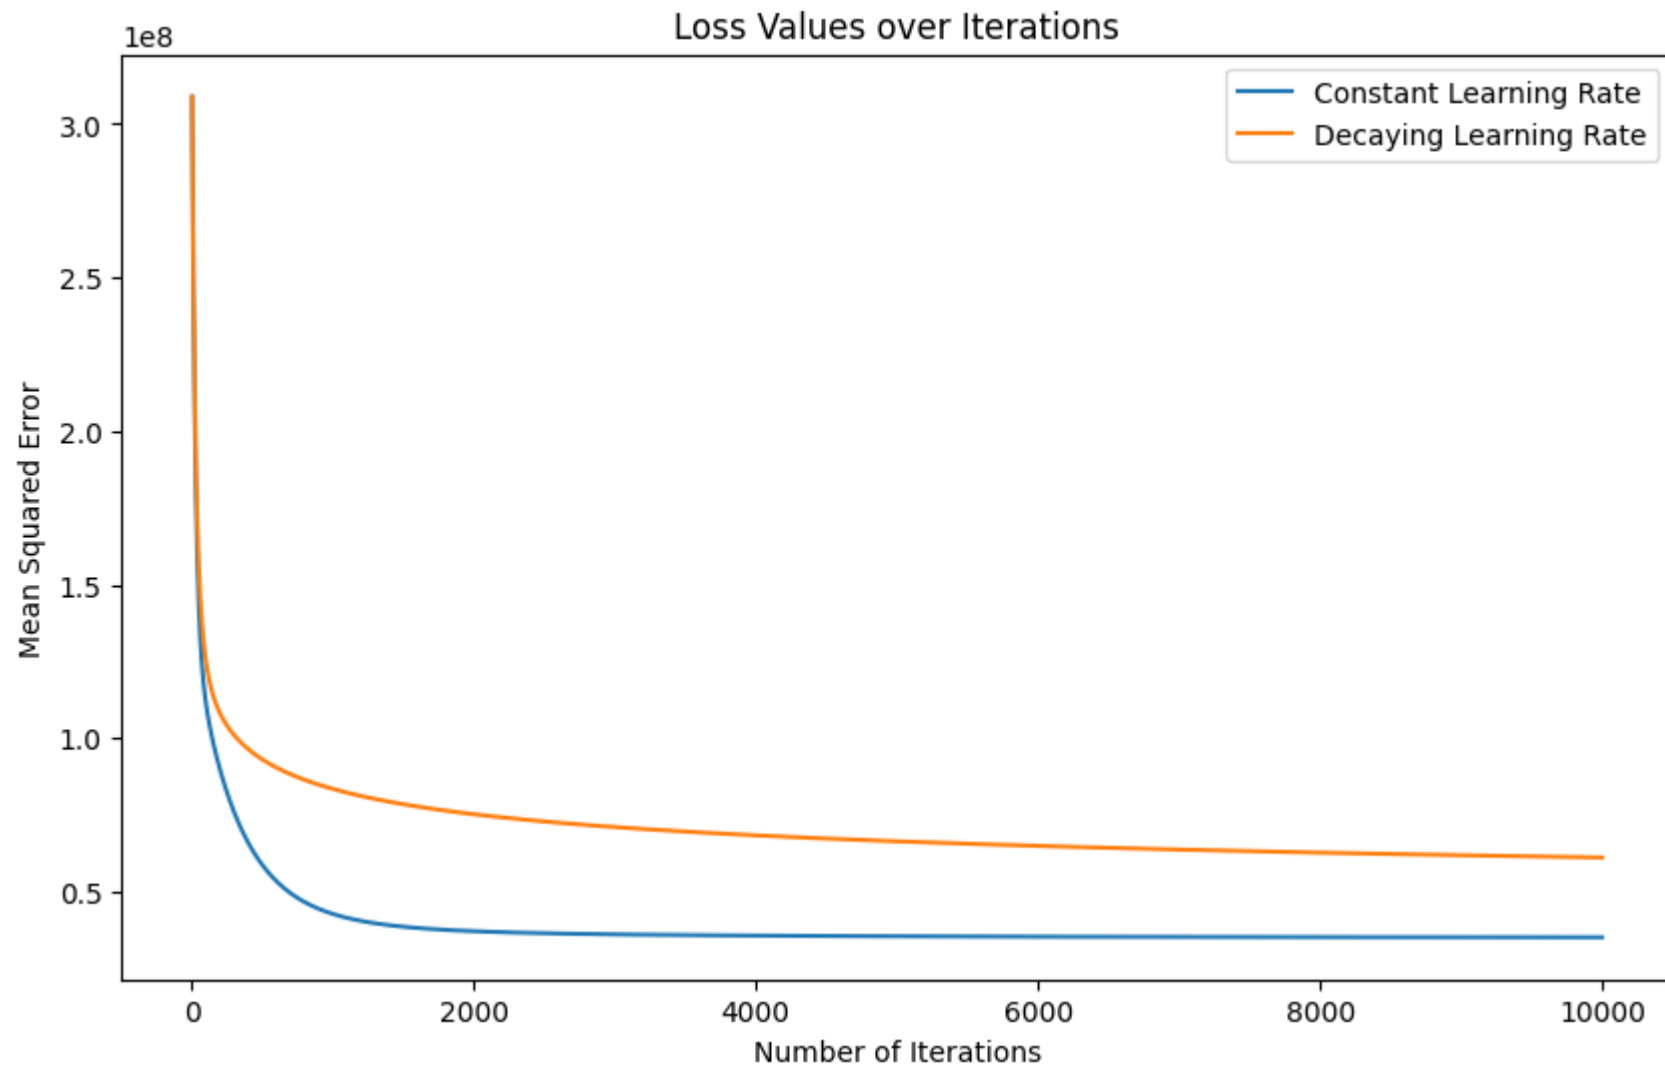
        predictions_decaying = X_b.dot(W_decaying)
        loss_values_decaying.append(mean_squared_error(y_train, predictions_decaying))

plt.figure(figsize=(10, 6))
plt.plot(range(1, n_iterations + 1), loss_values_constant, label='Constant Learning Rate')
plt.plot(range(1, n_iterations + 1), loss_values_decaying, label='Decaying Learning Rate')
plt.xlabel('Number of Iterations')
plt.ylabel('Mean Squared Error')
plt.title('Loss Values over Iterations')
plt.legend()
plt.show()

return W_constant, loss_values_constant, W_decaying, loss_values_decaying

W_constant, loss_values_constant, W_decaying, loss_values_decaying = gradient_descent(X_train, y_train, decay_rate=0.01)

```



## 4. Model Evaluation

```
In [40]: from sklearn.metrics import mean_absolute_error, mean_squared_error
def evaluate_model(X_test, y_test, W_constant, W_decaying):
    if len(X_test.shape) == 1: # For 1D Conversion
        X_test = X_test.reshape(-1, 1)
```

```

if len(y_test.shape) == 1: # For 1D Conversion since ytrain is series so using numpy here
    y_test = y_test.to_numpy().reshape(-1, 1)

X_test_b = np.c_[np.ones((X_test.shape[0], 1)), X_test]

predictions_test_constant = X_test_b.dot(W_constant)
predictions_test_decaying = X_test_b.dot(W_decaying)

# Computing MAE and MSE for both constant and decaying learning rates
mae_constant = mean_absolute_error(y_test, predictions_test_constant)
mse_constant = mean_squared_error(y_test, predictions_test_constant)

mae_decaying = mean_absolute_error(y_test, predictions_test_decaying)
mse_decaying = mean_squared_error(y_test, predictions_test_decaying)

# Histogram For Constant
error_distribution_constant = y_test - predictions_test_constant
plt.figure(figsize=(10, 6))
plt.hist(error_distribution_constant, bins=30, edgecolor='black')
plt.title('Error Distribution - Constant Learning Rate')
plt.xlabel('Error')
plt.ylabel('Frequency')
plt.show()

# Histogram For Decaying Rate
error_distribution_decaying = y_test - predictions_test_decaying
plt.figure(figsize=(10, 6))
plt.hist(error_distribution_decaying, bins=30, edgecolor='black')
plt.title('Error Distribution - Decaying Learning Rate')
plt.xlabel('Error')
plt.ylabel('Frequency')
plt.show()

print("Constant Learning Rate:")
print(f"MAE: {mae_constant:.2f}")
print(f"MSE: {mse_constant:.2f}")


print("\nDecaying Learning Rate:")
print(f"MAE: {mae_decaying:.2f}")

```

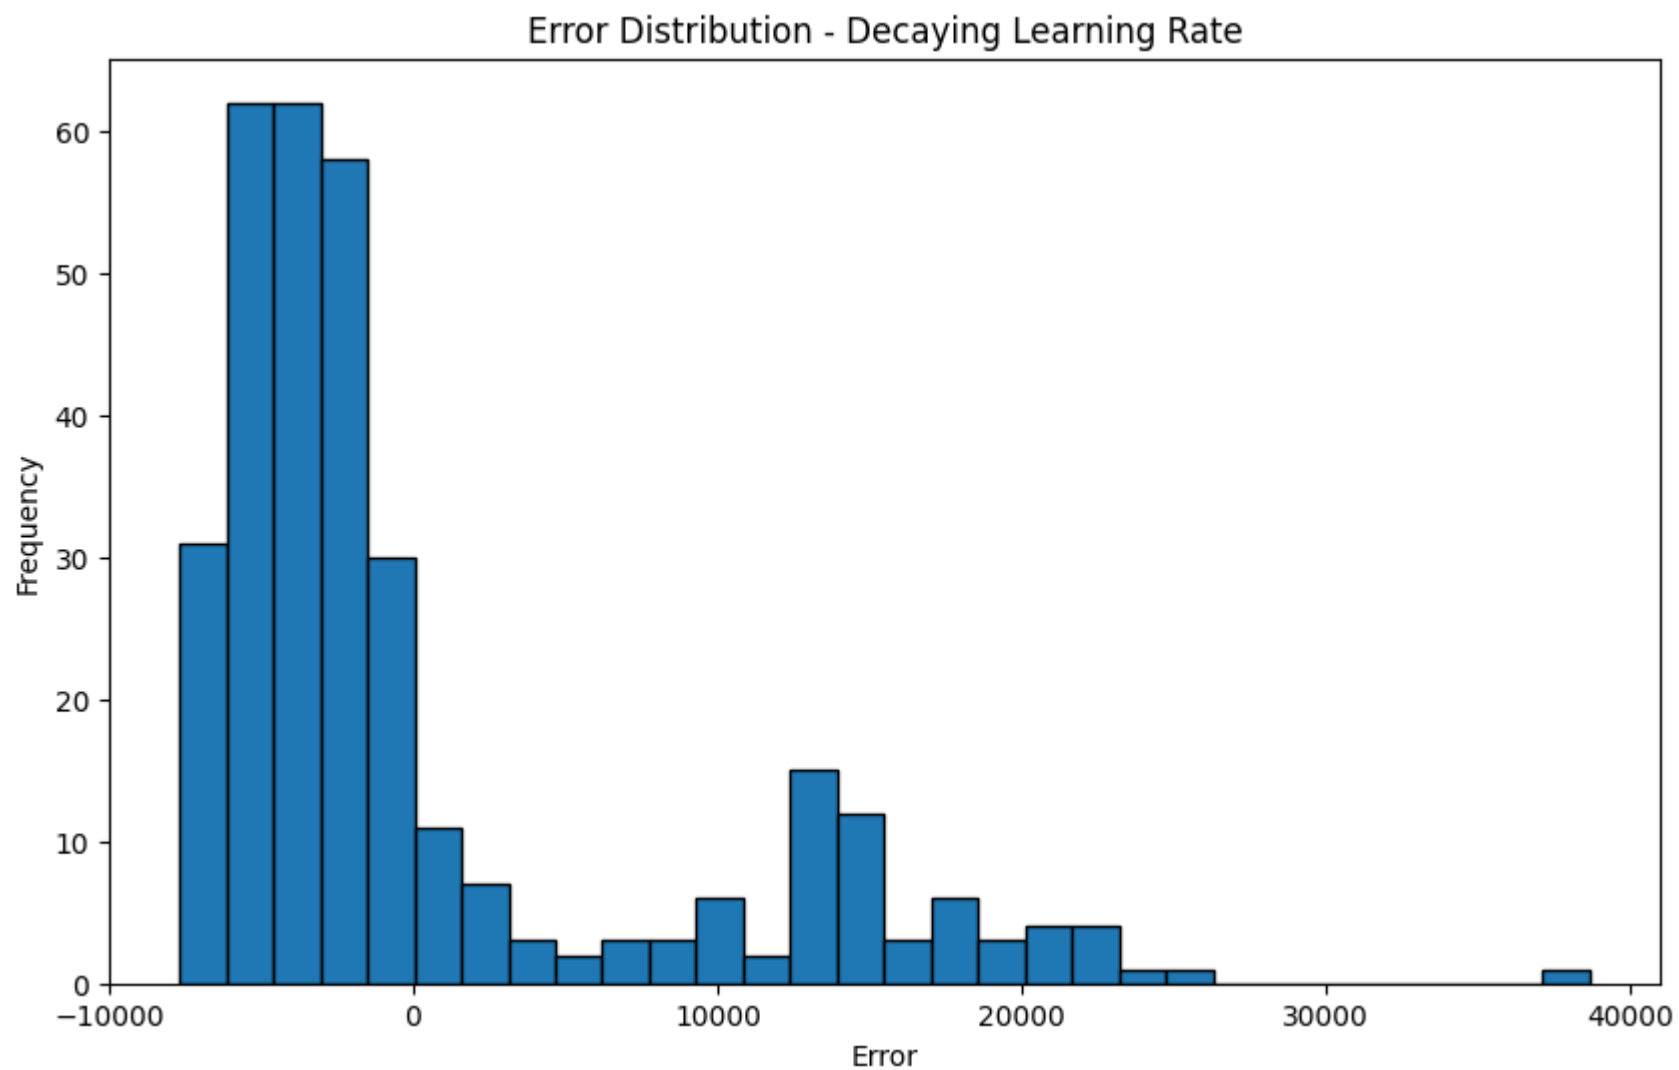
```
print(f"MSE: {mse_decaying:.2f}")
```

```
# Example usage
```

```
evaluate_model(X_test, y_test, W_constant, W_decaying)
```







Constant Learning Rate:

MAE: 4543.38

MSE: 43794979.02

Decaying Learning Rate:

MAE: 5988.32

MSE: 65776966.98

## 5. Learning Rate Analysis

```
In [41]: from sklearn.preprocessing import StandardScaler

def learning_rate_analysis(X, y, learning_rates, n_iterations=1000):

    # Convert y to a NumPy array if it's a Pandas Series
    if isinstance(y, pd.Series):
        y = y.to_numpy()

    # 2d to 1D Conversion
    if len(X.shape) == 1:
        X = X.reshape(-1, 1)
    if len(y.shape) == 1:
        y = y.reshape(-1, 1)

    X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

    scaler = StandardScaler()
    X_train_scaled = scaler.fit_transform(X_train)
    X_test_scaled = scaler.transform(X_test)

    W = np.random.randn(X_train_scaled.shape[1] + 1, 1)

    X_train_b = np.c_[np.ones((X_train_scaled.shape[0], 1)), X_train_scaled]

    loss_values = []

    for learning_rate in learning_rates:
        W_current = W.copy()

        local_loss_values = []

        for iteration in range(n_iterations):
            gradients = 1/X_train_b.shape[0] * X_train_b.T.dot(X_train_b.dot(W_current) - y_train)
            W_current = W_current - learning_rate * gradients

            predictions = X_train_b.dot(W_current)
```

```
        local_loss_values.append(mean_squared_error(y_train, predictions))

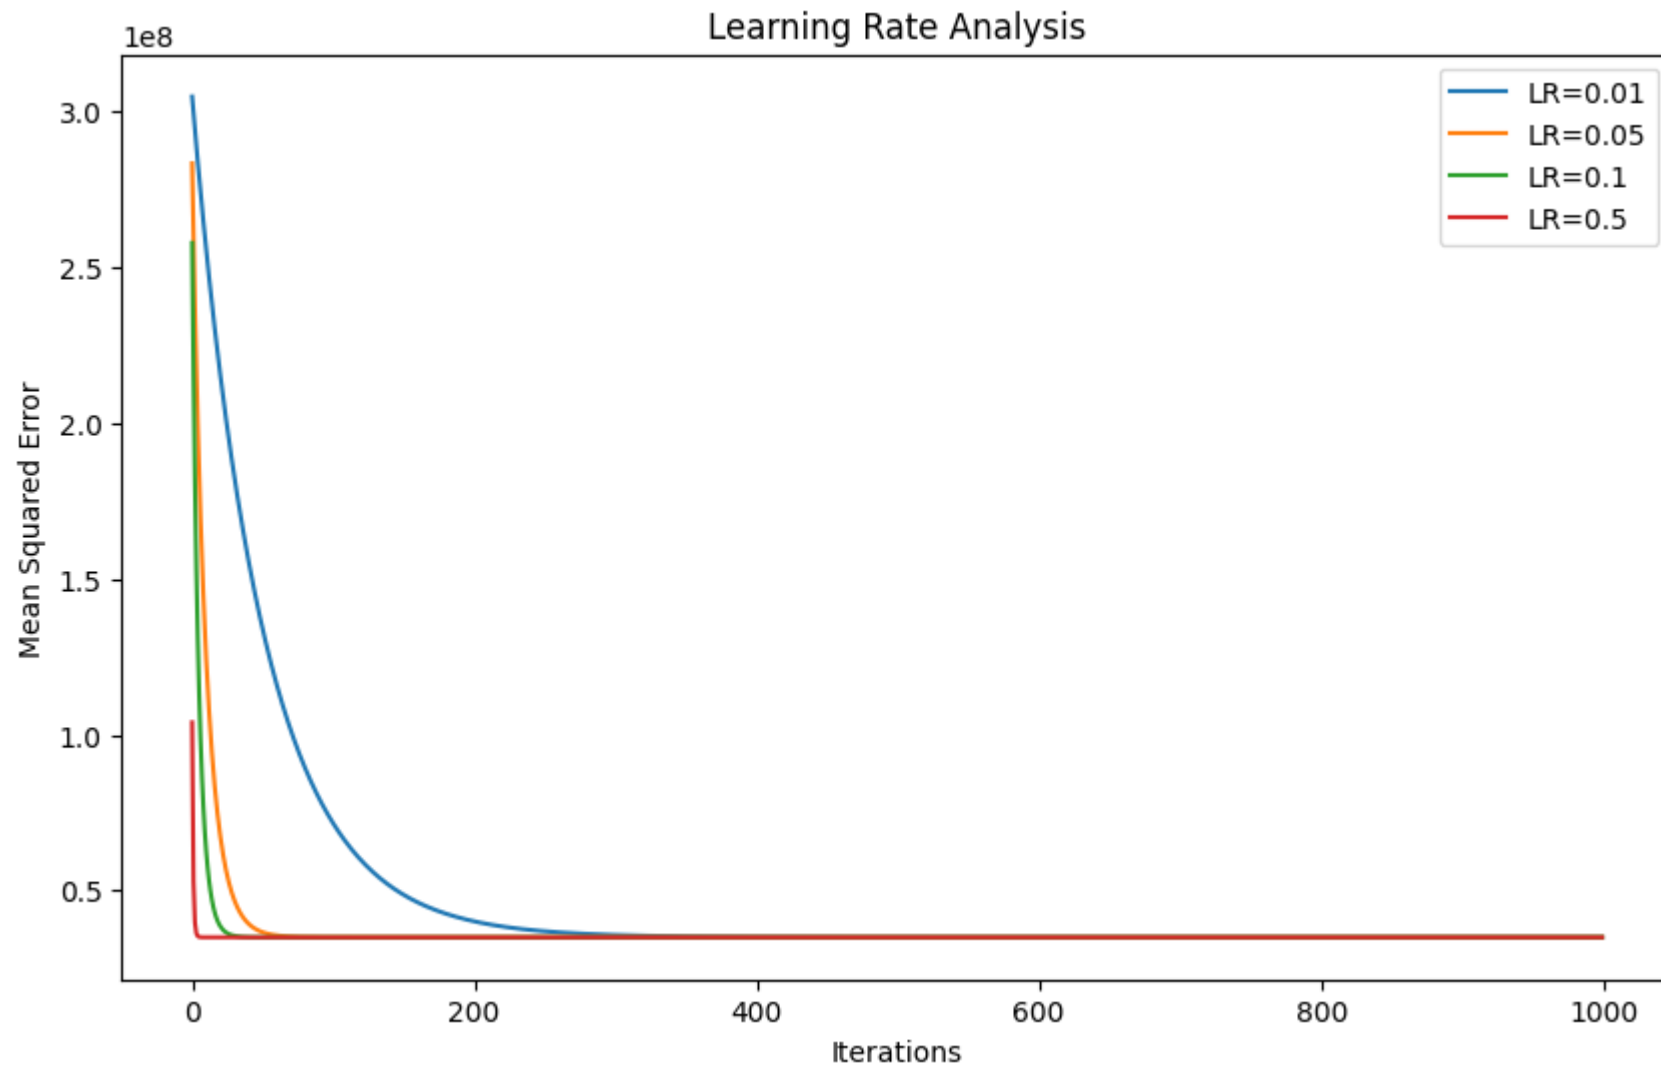
    loss_values.append(local_loss_values)

plt.figure(figsize=(10, 6))
for i, learning_rate in enumerate(learning_rates):
    plt.plot(range(n_iterations), loss_values[i], label=f'LR={learning_rate}')

plt.xlabel('Iterations')
plt.ylabel('Mean Squared Error')
plt.title('Learning Rate Analysis')
plt.legend()
plt.show()

return loss_values

learning_rates_to_analyze = [0.01, 0.05, 0.1, 0.5]
loss_values = learning_rate_analysis(X, Y, learning_rates_to_analyze)
```



## 6. Scikit-learn Implementation

```
In [42]: from sklearn.linear_model import LinearRegression
def scikit_learn_regression(X, y):
    X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```

scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

model = LinearRegression()

model.fit(X_train_scaled, y_train)

predictions = model.predict(X_test_scaled)

# MSE & MAE Comparision
mae = mean_absolute_error(y_test, predictions)
mse = mean_squared_error(y_test, predictions)

return mae, mse

mae, mse = scikit_learn_regression(X, Y)
print("MAE:", mae)
print("MSE:", mse)

```

MAE: 4743.839512329654

MSE: 47119179.47700916

## 7. Normal Equation Implementation

```

In [43]: def normal_equation_regression(X, y):
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

X_train_b = np.c_[np.ones((X_train_scaled.shape[0], 1)), X_train_scaled]

theta = np.linalg.inv(X_train_b.T.dot(X_train_b)).dot(X_train_b.T).dot(y_train)

X_test_b = np.c_[np.ones((X_test_scaled.shape[0], 1)), X_test_scaled]

predictions = X_test_b.dot(theta)

```

```

    mae = mean_absolute_error(y_test, predictions)
    mse = mean_squared_error(y_test, predictions)

    return mae, mse

mae_normal, mse_normal = normal_equation_regression(X, Y)
print("MAE (Normal Equation):", mae_normal)
print("MSE (Normal Equation):", mse_normal)

```

MAE (Normal Equation): 4743.839512329654

MSE (Normal Equation): 47119179.477009155

### Observation 💡💡💡 :

- MAE and MSE are almost Same for Linear Regression Model and Normal Equation because both use Linear Regression Model

```

In [44]: # MAE & MSE Calculation for Gradient_Descent
def gradient_descent(X, y, learning_rate=0.01, n_iterations=1000, decay_rate=None):
    if len(X.shape) == 1: # For 1D Conversion
        X = X.reshape(-1, 1)

    if len(y.shape) == 1: # For 1D Conversion since ytrain is series so using numpy here
        y = y.to_numpy().reshape(-1, 1)

    X_b = np.c_[np.ones((X.shape[0], 1)), X]

    m, n = X.shape[0], X.shape[1]

    W = np.random.randn(n + 1, 1)

    loss_values = []

    for iteration in range(n_iterations):
        gradients = 1/m * X_b.T.dot(X_b.dot(W) - y)

        if decay_rate:
            current_learning_rate = learning_rate / (1 + decay_rate * iteration)
        else:
            current_learning_rate = learning_rate

```

```

        W = W - current_learning_rate * gradients
        predictions = X_b.dot(W)
        loss_values.append(mean_squared_error(y, predictions))

    return W, loss_values

scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

W_gradient_descent, loss_values_gradient_descent = gradient_descent(X_train_scaled, y_train)
mae_gradient_descent = mean_absolute_error(y_test, np.c_[np.ones((X_test_scaled.shape[0], 1)), X_test_scaled].dot(W_gradient_d
mse_gradient_descent = mean_squared_error(y_test, np.c_[np.ones((X_test_scaled.shape[0], 1)), X_test_scaled].dot(W_gradient_de
print(f"Gradient Descent - MAE: {mae_gradient_descent:.4f}, MSE: {mse_gradient_descent:.4f}")

```

Gradient Descent - MAE: 4618.5731, MSE: 45290930.4499

```

In [45]: import time
         np.random.seed(42)
         X = np.random.rand(1000, 5)
         theta_true = np.random.rand(6, 1)
         X_b = np.c_[np.ones((1000, 1)), X]
         y = X_b.dot(theta_true) + np.random.randn(1000, 1) * 0.1 # Adding noise for realism
         X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

         scaler = StandardScaler()
         X_train_scaled = scaler.fit_transform(X_train)
         X_test_scaled = scaler.transform(X_test)

         # Computational efficiency for Gradient Descent
         start_time = time.time()
         W_gradient_descent, _ = gradient_descent(X_train_scaled, y_train)
         end_time = time.time()
         mae_gradient_descent = mean_absolute_error(y_test, np.c_[np.ones((X_test_scaled.shape[0], 1)), X_test_scaled].dot(W_gradient_d
         mse_gradient_descent = mean_squared_error(y_test, np.c_[np.ones((X_test_scaled.shape[0], 1)), X_test_scaled].dot(W_gradient_de
         ct_gradient = end_time-start_time
         print(f"Gradient Descent - Time: {ct_gradient:.4f}s, MAE: {mae_gradient_descent:.4f}, MSE: {mse_gradient_descent:.4f}")

         # For Normal Equation
         start_time = time.time()

```

```

mae_normal, mse_normal = normal_equation_regression(X_train, y_train)
end_time = time.time()
ct_normal = end_time - start_time
print(f"Normal Equation - Time: {ct_normal:.4f}s, MAE: {mae_normal:.4f}, MSE: {mse_normal:.4f}")

# For scikit-learn Linear Regression
start_time = time.time()
mae_scikit_learn, mse_scikit_learn = scikit_learn_regression(X_train, y_train)
end_time = time.time()
ct_scikit_learn = end_time - start_time
print(f"Scikit-learn Linear Regression - Time: {ct_scikit_learn:.4f}s, MAE: {mae_scikit_learn:.4f}, MSE: {mse_scikit_learn:.4f}")

```

Gradient Descent - Time: 0.7779s, MAE: 0.0897, MSE: 0.0120  
 Normal Equation - Time: 0.0000s, MAE: 0.0860, MSE: 0.0118  
 Scikit-learn Linear Regression - Time: 0.0105s, MAE: 0.0860, MSE: 0.0118

In [46]: mae\_gradient\_descent, mae\_normal, mae\_scikit\_learn

Out[46]: (0.089650517747645, 0.086048081464194, 0.08604808146419404)

In [47]: mse\_gradient\_descent, mse\_normal, mse\_scikit\_learn

Out[47]: (0.012018391740325483, 0.011786053686983744, 0.011786053686983738)

In [48]: ct\_gradient, ct\_normal, ct\_scikit\_learn

Out[48]: (0.7779483795166016, 0.0, 0.01051473617553711)

In [51]: **import** matplotlib.pyplot **as** plt

```

# Values
mae_values = [mae_gradient_descent, mae_normal, mae_scikit_learn]
mse_values = [mse_gradient_descent, mse_normal, mse_scikit_learn]
ct_values = [ct_gradient, ct_normal, ct_scikit_learn]

# Types
methods = ['Gradient Descent', 'Normal Equation', 'Scikit-learn']

# MAE plot
plt.figure(figsize=(10, 5))

```



```
plt.subplot(1, 3, 1)
plt.barh(methods, mae_values, color=['yellow', 'green', 'orange'])
plt.title('MAE Comparison')
plt.ylabel('MAE')

# MSE Plot
plt.subplot(1, 3, 2)
plt.barh(methods, mse_values, color=['yellow', 'green', 'orange'])
plt.title('MSE Comparison')
plt.ylabel('MSE')

# Computational Time Plot
plt.subplot(1, 3, 3)
plt.barh(methods, ct_values, color=['blue', 'green', 'orange'])
plt.title('Computational Time Comparison')
plt.ylabel('Computational Time (s)')

plt.tight_layout()
plt.show()
```

